

## Question 4

Let  $\{(Y_i, X_i)\}_{i=1}^n$  be an i.i.d. sequence of random vectors. Suppose that  $\mathbb{E}[X_i X_i']$  and  $\mathbb{E}[X_i Y_i]$  exists. Suppose further that there is no perfect colinearity in  $X_i$ , Hence  $\mathbb{E}[X_i X_i']$  is invertible.

**a**

Does it also follow that

$$\frac{1}{n} \sum_{i=1}^n X_i X_i'$$

is invertible?

No. As a trivial case, consider when  $n = 1, k = 2$  and  $X_2 \sim \mathcal{N}(1, 1)$ . Let  $a$  be any realization of  $X_2$ .

$$\frac{1}{n} \sum_{i=1}^n X_i X_i' = (1, a)'(1, a) = \begin{pmatrix} 1 & a \\ a & a^2 \end{pmatrix}$$

We can see that the second column is  $a$  times the first column, and the matrix is not invertible. This occurs because for any vector  $x \in \mathbb{R}^k$ ,  $xx'$  always has rank 1.

**b**

For any  $\lambda_n > 0$  show that

$$\frac{1}{n} \sum_{i=1}^n (X_i X_i' + \lambda_n \mathbb{I})$$

is invertible.

Note that this can be rewritten as

$$\left[ \frac{1}{n} \sum_{i=1}^n X_i X_i' \right] + \lambda_n \mathbb{I}$$

For any given  $i$ ,  $X_i X_i'$  is positive semi-definite. The sum of positive semi-definite matrices is also positive semi-definite. This tells us that the first matrix is always positive semi-definite.

$$\frac{1}{n} \sum_{i=1}^n X_i X_i' \succeq 0$$

It is obvious that  $\lambda_n \mathbb{I}$  is a positive definite matrix. The sum of a positive definite matrix and a positive semi-definite matrix is positive definite.

Proof: Let  $A$  be a positive semi-definite matrix, and  $B$  be a positive definite matrix. Then  $\forall x \in \mathbb{R}^k$ ,  $x' B x > 0$  and  $x' A x \geq 0$ . Consider two cases:

Case 1:  $x \in \mathbb{R}^k, x'Ax > 0, x'Bx > 0$ . Then:

$$\begin{aligned}(x'A + x'B)X &> 0 \\ x'(A + B)x &> 0\end{aligned}$$

Case 2:  $x \in \mathbb{R}^k, x'Ax = 0, x'Bx > 0$  Then:

$$\begin{aligned}x'Ax + x'Bx &> 0 \\ (x'A + x'B)X &> 0 \\ x'(A + B)x &> 0\end{aligned}$$

This tells us that:

$$\left[ \frac{1}{n} \sum_{i=1}^n X_i X_i' \right] + \lambda_n \mathbb{I} \succ 0$$

Any positive definite matrix has strictly positive eigenvalues, and therefore has a strictly positive determinant. This implies that the matrix is invertible.

**c**

Suppose that  $\lambda_n \rightarrow 0$  as  $n \rightarrow \infty$ . Find the limit in probability of

$$\tilde{\beta}_n = \left( \frac{1}{n} \sum_{i=1}^n (X_i X_i' + \lambda_n \mathbb{I}) \right)^{-1} \left( \frac{1}{n} \sum_{i=1}^n X_i Y_i \right)$$

From the weak law of large numbers, we know that  $\frac{1}{n} \sum_{i=1}^n X_i X_i' \xrightarrow{p} \mathbb{E}[X X']$  and  $\frac{1}{n} \sum_{i=1}^n X_i Y_i \xrightarrow{p} \mathbb{E}[XY]$ .

We wish to show that

$$\frac{1}{n} \sum_{i=1}^n X_i X_i' + \lambda_n \mathbb{I} \xrightarrow{p} \frac{1}{n} \sum_{i=1}^n X_i X_i'$$

Applying the definition of convergence in probability.

$$\lim_{n \rightarrow \infty} \Pr \left( \left| \frac{1}{n} \sum_{i=1}^n X_i X_i' + \lambda_n \mathbb{I} - \frac{1}{n} \sum_{i=1}^n X_i X_i' \right| < \epsilon \right) = \lim_{n \rightarrow \infty} \Pr(|\lambda_n \mathbb{I}| < \epsilon)$$

We will consider this on an element-wise basis. Note that if we are not on a diagonal,  $(\lambda_n \mathbb{I})_{ij} = 0$ . So we may restrict ourselves to the diagonal elements of this matrix. However all the diagonal elements are the same, so this question amounts to the convergence of  $|\lambda_n|$ . Since  $\lambda_n$  is non-random:

$$\lim_{n \rightarrow \infty} \Pr(|\lambda_n| < \epsilon) = 1$$

As we have assumed that  $\lambda_n \rightarrow 0$  above.

Thus

$$\frac{1}{n} \sum_{i=1}^n X_i X_i' + \lambda_n \mathbb{I} \xrightarrow{p} \frac{1}{n} \sum_{i=1}^n X_i X_i' \xrightarrow{p} \mathbb{E}[X X']$$

As multiplication and inverting a matrix are continuous functions, we may apply the continuous mapping theorem to get that

$$\tilde{\beta}_n = \left( \frac{1}{n} \sum_{i=1}^n (X_i X_i' + \lambda_n \mathbb{I}) \right)^{-1} \left( \frac{1}{n} \sum_{i=1}^n X_i Y_i \right) \xrightarrow{p} \mathbb{E}[X X']^{-1} \mathbb{E}[X Y] = \beta$$

## Question 8

**a**

---

```

1 data <- read.csv( "ps4.csv" )
2
3 k <- ncol(data)
4 N <- nrow(data)
5
6 ## Since we are not calling lm, we want to do matrix algebra, we need
7 ## R to not store this stuff as a data frame. What a terrible language.
8
9 Y <- as.matrix(data$y)
10 X <- as.matrix(cbind( rep(1,N), data[,2:3] ))
11
12 ## Remember that matrix multiplication uses the %*%
13 mat <- t(X)%*%X
14
15 ## Rather than using inverses, let's be numerically stable and use the
16 ## Cholesky decomp and forward/back substitution for legitimate answers
17 F <- chol(mat)
18
19 ## We now have X'Xβ = X'Y
20 ## This is equivalent to F'Fβ = X'Y
21 ## Thus β = F-1F'-1X'Y
22
23 ## Note that F' is lower triangular so we use forward substitution.
24 beta <- backsolve( F, forwardsolve( t(F), t(X)%*%Y ) )

```

---

Our estimated values of  $\beta$  are: (0.1680066, 1.0843565, 0.9203671)'.

---

**b**


---

```

25 ## Now lets build our variance estimates.
26
27 outerproduct <- function( row ){
28     row%*%t(row)
29 }
30
31 ## We are interested in estimating  $(\frac{1}{n} \sum_{i=1}^n X_i X_i')^{-1}$ 
32
33
34 ## The inner apply() forms the outer product matrices, the outer
35 ## averages over them The matrix() reforms them as a matrix since
36 ## apply flattens them. This is equivalent to just doing
37 ##  $\frac{1}{n} X'X$ , I just wanted some R practice.
38 outerProductGradient <- matrix( apply( apply( X, 1, outerproduct ), 1,
39                                     mean ), nrow = k, ncol = k )
40
41 ## Mama told me to never invert a matrix on a computer
42 varF <- chol( outerProductGradient )
43 informationEstimate <- backsolve( varF, forwardsolve( t(varF), diag(k) ) )
44
45 ## Now lets get the heteroskedasticity-robust version of this bad boy.
46 ## We multiply the matrix of  $X_i X_i'$  by  $\hat{u}_i^2$  component wise, hence no %
47 monstronsity <- matrix( apply(
48     matrix( rep( (Y - X%*%beta)^2, k*k ), nrow=k*k, ncol = N, byrow = TRUE )
49     * apply( X, 1, outerproduct ), 1, mean ), nrow = k, ncol = k )
50
51 ## This is what are interested in:  $\mathbb{V}(\hat{\beta}_N|X)$ 
52 condVarHetero <- informationEstimate%*%monstronsity%*%informationEstimate
53
54
55 ## Note that it's possible to just use matrix operations to get there
56 ## I just chose this way for practice and to have it look like the notes.
57 ## One could always do  $(X'X)^{-1} X' \hat{\Sigma}_N X (X'X)^{-1}$ 

```

---

Our estimated Variance-Covariance Matrix of  $\hat{\beta}_N$  is:

$$\mathbb{V}(\hat{\beta}_N|X) = \begin{pmatrix} 4.8905355 & 0.4493318 & -1.6478739 \\ 0.4493318 & 0.4517238 & -0.3702895 \\ -1.6478739 & -0.3702895 & 0.7567006 \end{pmatrix}$$

## c

---

```

58 ## Now we face multiple linear restrictions in the form of  $R\beta = r$ 
59
60 ## We don't really know anything about the nature of  $R\mathbb{V}(\hat{\beta}_N)$ 
61 ## So we can't rely on any decompositions, and we'll let solve() work here
62 multipleLinearTest <- function( R, r, N, beta, Var ){
63     N*t(R%*%beta - r )%*%solve(R%*%Var%*%t(R))%*%(R%*%beta -r )
64 }
65
66
67 R <- matrix( c( 0, 0, 1, 0 ,0,1 ), nrow = 2, ncol = 3 )
68 r <- c( 1, 1 )
69
70 ## This is free to be changed.
71 alpha <- .05
72
73 ## This c is the critical value used in a hypothesis test
74 c <- qchisq( alpha, df = 2, lower.tail = FALSE )
75
76
77 testStat <- multipleLinearTest( R, r, N, beta, condVarHetero )
78 pValue <- pchisq( testStat, df = 2, lower.tail = FALSE )

```

---

Our test statistic value is 1.599558 and our p-value is: 0.4494283

## d

---

```

79 ## Testing:  $f(\beta) = (\beta_1 - \beta_2)^2 = 0$ 
80 ## However we need the rows of the total derivative to be linearly
81   ↪ independent.
82 ##  $\nabla f(\beta) = (0, 2(\beta_1 - \beta_2), -2(\beta_1 - \beta_2))'$ 
83 ## The rows are not linearly independent - The standard nonlinear test
84   ↪ will not work.
85
86 ## Worse yet, if we attempt to simply take the square root of both
87 ## sides we lose the reliability as this is a Wald-Test. Wald Tests
88 ## are not invariant to non-linear Transforms. This means we want to
89 ## use a likelihood-ratio test, which is. However if we do not want to
90 ## assume normality of Y and then the GLM framework to get a
91 ## likelihood-ratio test, we can just stand for the errors in the Wald
92   ↪ Test.

```

---

```
91 ## Our test is simply testing if  $\beta_1 - \beta_2 = 0$ 
92
93 R <- matrix( c( 0, 1, -1 ), nrow = 1, ncol = 3 )
94 r <- c(0)
95
96 ## I just copy and pasted the previous code
97 c <- qchisq( alpha, df = 1, lower.tail = FALSE )
98 testStat <- fischerFTest( R, r, N, beta, condVarHetero )
99 pValue <- pchisq( testStat, df = 1, lower.tail = FALSE )
```

---

Our test statistic value is 1.379809 and our p-value is: 0.2401337