

Optimization Conscious Econometrics Pset 2

Part 2

Timothy Schwieg

1 Write the likelihood for model (1) with both logit and linear links

1.1 Linear Link

For the linear link, we note that $x_{i,t}\beta$ is a constant, and that $\alpha_i \sim \mathcal{N}(0, \tau^2)$ so $x_{i,t}\beta + \alpha_i \sim \mathcal{N}(x_{i,t}\beta, \tau^2)$.

However, for each observation for person i , there is the same draw of α_i . The object in question is $x_i\beta + 1_T\alpha_i$ where 1_T is a T -dimensional vector containing all ones, and x_i is the vector containing all T measurements from consumer i .

The joint distribution of this object is given by:

$$x_i\beta + 1_T\alpha_i \sim \mathcal{N}(x_i\beta, \tau^2 1_T 1_T')$$

Letting $\Sigma = \tau^2 1_T 1_T'$

The Likelihood function is simply the product of the densities evaluated at their outcomes.

$$L = \prod_{i=1}^n \frac{1}{\sqrt{2\pi |\Sigma|}} \exp\left(-\frac{1}{2}(y_i - x_i)' \Sigma^{-1} (y_i - x_i)\right)$$

1.2 Logit Link Function

I invert the logit link function so that we reach the model where $y_{i,t} = g(x_{i,t}\beta + \alpha_i)$ where

$$g(y) = \begin{pmatrix} \frac{\exp(y_1)}{1+\exp(y_1)} \\ \dots \\ \frac{\exp(y_t)}{1+\exp(y_t)} \\ \dots \\ \frac{\exp(y_T)}{1+\exp(y_T)} \end{pmatrix}$$

This is a bijection from $\mathbb{R}^T \rightarrow (0, 1)^T$. As such we can define the inverse mapping g^{-1} .

$$g^{-1}(y) = \begin{pmatrix} \log\left(\frac{y_1}{1-y_1}\right) \\ \dots \\ \log\left(\frac{y_t}{1-y_t}\right) \\ \dots \\ \log\left(\frac{y_T}{1-y_T}\right) \end{pmatrix}$$

Note that $\frac{\exp(y_i)}{1+\exp(y_i)} = 1 - \frac{1}{1+\exp(y_i)}$. Its derivative is $g(x)(1 - g(x))$. The Jacobian of the transformation $g(x)$ is given by:

$$J = \begin{pmatrix} g(y_1)(1 - g(y_1)) & 0 & \dots & 0 \\ 0 & g(y_2)(1 - g(y_2)) & \dots & 0 \\ \dots & \dots & \dots & \dots \\ 0 & 0 & \dots & g(y_T)(1 - g(y_T)) \end{pmatrix}$$

The determinant of the Jacobian matrix is then:

$$|J| = \prod_{t=1}^T g(y_t)(1 - g(y_t))$$

The distribution of the transformation is then given by:

Let g_i denote the vector of $g(y_i)$.

$$\begin{aligned} f_g &= \frac{f_{y_i}(y_i)}{|J|} \\ &= \frac{f_{y_i}(g^{-1}(g_i))}{|J|} \end{aligned}$$

Using the results from the previous section, and remembering that $\Sigma = \tau^2 1_T 1_T'$. The likelihood function can be calculated to be:

$$L = \prod_{i=1}^N \frac{1}{\prod_{t=1}^T g_{i,t}(1 - g_{i,t})} \frac{1}{\sqrt{2\pi} |\Sigma|} \exp\left(-\frac{1}{2}(g^{-1}(g_i) - x_i\beta)' \Sigma^{-1} (g^{-1}(g_i) - x_i\beta)\right)$$

2 b

In the case of the linear link, i.e., model (2) show that

$$\int f(y|\beta, \alpha) f(\alpha|\tau^2) d\alpha \propto \mathcal{N}(x\beta, \sigma^2 + \tau^2)$$

The likelihood for y conditional on β, α is normally distributed.

$$f(y|\beta, \alpha) = \prod_{i=1}^I \prod_{t=1}^T \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y_{i,t} - x_{i,t}\beta - \alpha_i)^2}{2\sigma^2}\right)$$

We may also note that α_i is also normally distributed when conditioned upon τ^2 . Therefore:

$$f(\alpha|\tau^2) = \prod_{i=1}^I \frac{1}{\sqrt{2\pi\tau^2}} \exp(-\frac{\alpha_i^2}{2\tau^2})$$

Taking the integral of the product of these two terms:

$$\int \prod_{t=1}^T \prod_{i=1}^I \frac{1}{\sqrt{2\pi(\sigma^2\tau^2)}} \exp(-\frac{\alpha_i^2\sigma^2 + \tau^2(y_{i,t} - x_{i,t}\beta - \alpha_i)^2}{2\tau^2\sigma^2}) d\alpha$$

Which we are given is equal to the integrated likelihood function $L(y|\beta, \tau)$.

This tells us that the distribution of $\int f(y|\alpha, \beta, \tau)f(\alpha|\tau)$ is the same as the distribution of y unconditional on any knowledge of α beyond its variance and β . If we did not observe α_i , then the unobserved terms in model (2) are $\alpha_i \sim \mathcal{N}(\tau^2)$ and $\epsilon_{i,t} \sim \mathcal{N}(\sigma^2)$ so therefore the sum of the two is distributed $\mathcal{N}(0, \sigma^2 + \tau^2)$. Adding in the constant term arrive at the distribution of a single $y_{i,t} \sim \mathcal{N}(x_{i,t}\beta, \sigma^2 + \tau^2)$.

From this we can conclude:

$$\int f(y|\beta, \alpha)f(\alpha|\tau^2)d\alpha = L(y|\beta, \tau) \propto \mathcal{N}(x\beta, \sigma^2 + \tau^2)$$

3 c

For an S large enough, we know that $\frac{1}{S} \sum_{s=1}^S f(y|\beta^* \sigma_s^*) \xrightarrow{p} L(y|\beta, \tau)$.

This guess is an unbiased estimate of the likelihood, and using pseudo-marginal MCMC will results in a stationary distribution that is sampling from the likelihood correctly. After we have continued through the burnout period, and if the distribution is mixing well, we will have a sample of $\beta, \tau, \alpha_{(s)}$ where each element in $\alpha_{(s)}$ is distributed normally with variance of τ^2 . Since α is independent of all of the other parameters in the distribution, its marginal distribution is just a normal distribution with variance of τ^2 . Thus we are sampling from the correct distribution for α .

4 d

Algorithm A is implemented as follows:

```

1  function StandardMCMC( startVal, N::Int64)
2       $\beta$  = Vector{Float64}(undef, N+1)
3       $\tau$  = Vector{Float64}(undef, N+1)
4       $\alpha$  = Vector{Float64}(undef, N+1)
5      coinFlips = rand(Uniform(), N)
6       $\beta$ [1] = startVal[1]
7       $\tau$ [1] = startVal[2]
8       $\alpha$ [1] = startVal[3]
9      for i in 1:N
10         betaStar = betaSimFun(  $\beta$ [i])
11         tauStar = tauSimFun( $\tau$ [i])#qSimFun(  $\tau$ [i], qArgs )
12         alphaStar = rand(Normal(0, tauStar), 1)[1]
13
14         p1 = pFunStandard( betaStar, tauStar, alphaStar, Y, X, 1.0 )

```

```

15     # p1 = pFunLogit( betaStar, tauStar, alphaStar, Y,X )
16     p2 = qFun( $\beta[i]$ ,  $\tau[i]$ , betaStar, tauStar)
17     p3 = pFunStandard(  $\beta[i]$ ,  $\tau[i]$ , a[i], Y,X, 1.0 )
18     # p3 = pFunLogit(  $\beta[i]$ ,  $\tau[i]$ , a[i], Y,X )
19     p4 = qFun(betaStar, tauStar,  $\beta[i]$ ,  $\tau[i]$ )
20
21
22     if( coinFlips[i] <= p1*p2 / (p3*p4))
23          $\beta[i+1]$  = betaStar
24          $\tau[i+1]$  = tauStar
25         a[i+1] = alphaStar
26     else
27          $\beta[i+1]$  =  $\beta[i]$ 
28          $\tau[i+1]$  =  $\tau[i]$ 
29         a[i+1] = a[i]
30     end
31 end
32 return [ $\beta$ ,  $\tau$ , a]
33 end

```

Algorithm B is implemented as follows:

```

34 function MultiplyMCMC( startVal, N::Int64, S::Int64 )
35      $\beta$  = Vector{Float64}(undef,N+1)
36      $\tau$  = Vector{Float64}(undef,N+1)
37     a = Vector{Vector{Float64}}(undef,N+1)
38     coinFlips = rand(Uniform(),N)
39      $\beta[1]$  = startVal[1]
40      $\tau[1]$  = startVal[2]
41     a[1] = ones(S)*startVal[3]#rand(Normal(0, $\tau[1]$ ),S)
42     for i in 1:N
43         betaStar = betaSimFun(  $\beta[i]$ )
44         tauStar = tauSimFun( $\tau[i]$ )#qSimFun(  $\tau[i]$ , qArgs )
45         alphaStar = rand(Normal(0,tauStar),S)
46
47         p1 = exp(sum( log(pFunStandard( betaStar, tauStar, alphaStarZ, Y,X, 1.0 )) for alphaStarZ in alphaStar
48         ↪ ))
49         p2 = qFun( $\beta[i]$ ,  $\tau[i]$ , betaStar, tauStar)
50         p3 = exp(sum( log(pFunStandard(  $\beta[i]$ ,  $\tau[i]$ , alphaStarZ, Y,X, 1.0 )) for alphaStarZ in a[i] ))
51         p4 = qFun(betaStar, tauStar,  $\beta[i]$ ,  $\tau[i]$ )
52
53         quant = p1*p2/(p3*p4)
54
55         if( coinFlips[i] <= p1*p2/(p3*p4))
56              $\beta[i+1]$  = betaStar
57              $\tau[i+1]$  = tauStar
58             a[i+1] = alphaStar
59         else
60              $\beta[i+1]$  =  $\beta[i]$ 
61              $\tau[i+1]$  =  $\tau[i]$ 
62             a[i+1] = a[i]
63         end
64     end
65     return [ $\beta$ ,  $\tau$ , a]
66 end

```

I implement algorithm C as:

```

66 function PseudoMarginalMCMC( startVal, N::Int64, S::Int64, qFun, pFun, pArgs)
67      $\beta$  = Vector{Float64}(undef,N+1)
68      $\tau$  = Vector{Float64}(undef,N+1)
69     a = Vector{Vector{Float64}}(undef,N+1)
70     coinFlips = rand(Uniform(),N)

```

```

71     β[1] = startVal[1]
72     τ[1] = startVal[2]
73     α[1] = ones(S)*startVal[3]#rand(Normal(0,τ[1]),S)
74     for i in 1:N
75         betaStar = betaSimFun( β[i])
76         tauStar = tauSimFun(τ[i])#qSimFun( τ[i], qArgs )
77         alphaStar = rand(Normal(0,tauStar),S)
78
79         p1 = mean(pFunStandard( betaStar, tauStar, alphaStarZ, Y,X, σ² ) for alphaStarZ in alphaStar)
80         #p1 = mean(pFunLogit( betaStar, tauStar, alphaStarZ, Y,X ) for alphaStarZ in alphaStar)
81         #println(p1)
82         p2 = qFun(β[i], τ[i], betaStar, tauStar)
83         p3 = mean(pFunStandard(β[i], τ[i], alphaStarZ, Y,X, σ² ) for alphaStarZ in α[i])
84         #p3 =mean(pFunLogit(β[i], τ[i], alphaStarZ, Y,X ) for alphaStarZ in α[i])
85         p4 = qFun(betaStar, tauStar, β[i], τ[i])
86
87         quant = p1*p2/(p3*p4)
88
89         if( coinFlips[i] <= p1*p2/(p3*p4))
90             β[i+1] = betaStar
91             τ[i+1] = tauStar
92             α[i+1] = alphaStar
93         else
94             β[i+1] = β[i]
95             τ[i+1] = τ[i]
96             α[i+1] = α[i]
97         end
98     end
99
100     return [β, τ, α]
101 end

```

The functions used to simulate the data are as follows:

```

102 function pFunStandard( β, τ, α, y, x, σ² ) #, τ² )
103
104     betaPrior = -.5*log( 2*pi*priorMuSigma) - .5*(priorMuBeta - β)^2/ priorMuSigma
105
106     tauPrior =-.5*log( 2*pi*priorTauSigma) - .5*(priorTauBeta - β)^2/ priorTauSigma
107
108     likelihood = sum(-.5*log( 2*pi*σ²) - .5*(y[j,i] - x[j,i]*β - α)^2/ σ² for (i,j) in zip(1:I,1:T))
109
110     return exp(likelihood)#+betaPrior+tauPrior)#1.0 / sqrt( 2 * pi * σ²) * exp( -.5*(y - x*β)^2 / σ²)
111     #return alphaDensity*yDensity
112 end
113
114 function logit( x::Float64 )
115     return exp(x) / ( 1.0 + exp(x))
116 end
117
118 function logitMinus( x::Float64)
119     return 1.0 / (1.0+exp(x))
120 end
121
122
123
124 function pFunLogit( β, τ, α, y, x)
125     # Now lets consider when  $Y \sim \text{Bernoulli}(\ell(x\beta + \alpha_i))$ 
126     # Where  $\ell(x) = \frac{\exp(x)}{1+\exp(x)}$ 
127     # So the likelihood function is given by the pdf:  $\ell(x)^y(1 - \ell(x))^{(1-y)}$ 
128     # Log likelihood is:  $y \log \ell(x) + (1 - y) \log(1 - \ell(x))$ 
129
130     betaPrior = -.5*log( 2*pi*priorMuSigma) - .5*(priorMuBeta - β)^2/ priorMuSigma
131     #(1/sqrt(2*pi*priorMuSigma))*exp( (priorMuBeta-β)^2 / (2*priorMuSigma) )
132     tauPrior =-.5*log( 2*pi*priorTauSigma) - .5*(priorTauBeta - β)^2/ priorTauSigma

```

```

133      #(1/sqrt(2*pi*priorTauSigma))*exp( (priorTauBeta- $\tau$ )^2 / (2*priorTauSigma) )
134      likelihood = sum( y[j,i]*log(logit(x[j,i]* $\beta$  + a)) +
135      (1-y[j,i])*log(logitMinus(x[j,i]* $\beta$  + a)) for (i,j) in zip(1:I,1:T))
136      return exp(likelihood+betaPrior+tauPrior)
137  end
138
139
140  function betaFun( beta, betaOld )
141      return 1 / (2*.1)
142  end
143
144  function betaSimFun( xOld )
145      return rand(Uniform(xOld-.1,xOld+.1),1)[1]
146  end
147
148
149  function tauFun( tauNew, tauOld )
150      return 1.0 / ( min(tauOld+.01,5.0)-max(tauOld-.01,0.025))
151  end
152
153  function tauSimFun( tauOld)
154      return rand(Uniform(max(tauOld-.01,.025),min(tauOld+.01,5.0)),1)[1]
155  end
156
157  function qFun( beta, tau, betaOld, tauOld)
158      return betaFun( beta, betaOld)*tauFun(tau,tauOld)
159  end

```

I find that for the linear case, algorithm A and C converge to nearly the same value, with algorithm A doing it in much less computations. For even modest values of S however, I found that Algorithm B encountered numerical problems that resulted in an acceptance probability of zero. This likely occurs as we are multiplying many tiny probabilities together, particularly for initial prior values which have some distance from the true distribution. Even using double precision floating point numbers, for extremely small numbers this cannot be resolved easily.

I attempted several simple numerical techniques to avoid this issue, including changing products to the exponent of the sum of logarithms, but I was unable to resolve the acceptance probability defaulting to 0. This problem resulted from extremely low likelihoods even for small amounts of data. The product of these low probabilities (on the order of magnitude of -60) cannot be resolved as non-zero even for trivial values of S such as $S = 5$.

Between the two remaining algorithms: A and C, I found that after a large sample $N = 10000$, $S = 1000$ where there were 1000 individuals with 100 data points each, both converged to the same result. For the linear model, I gave a normal prior to both β and τ , on simulated data.

For the pseudo-marginal MCMC simulation, I found that there was a 35% acceptance probability for the distribution, and the auto-correlation function is shown as well.

For the single α sample (algorithm A), I found that there was a significantly lower acceptance probability (5%). This led to a much poorer mixture. However, the mixture under the pseudo-marginal MCMC is still not ideal, as there is quite a bit of noise in the distribution of τ despite the prior being a normal distribution, and the data being simulated from a normal distribution.

It is likely that this problem could be resolved by a better choice of sampling than the plain Monte-Carlo integration employed in the suggestion. If the points were suggested based

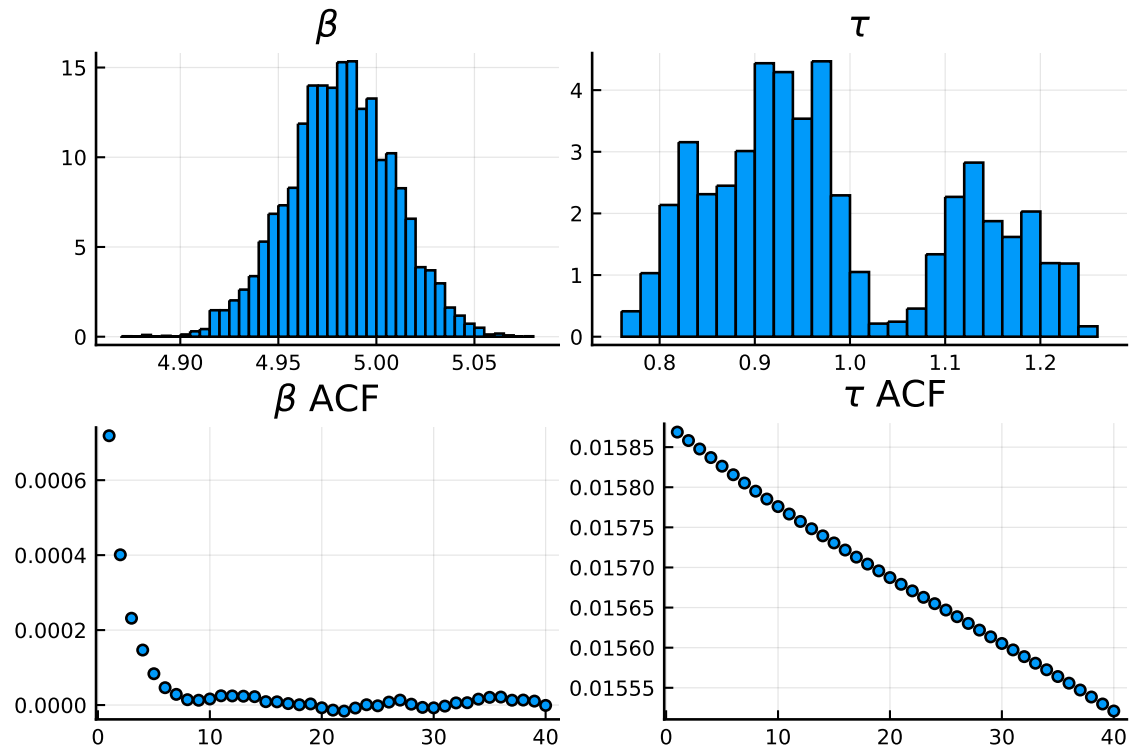


Fig. 1: Pseudo-Marginal MCMC Linear Link

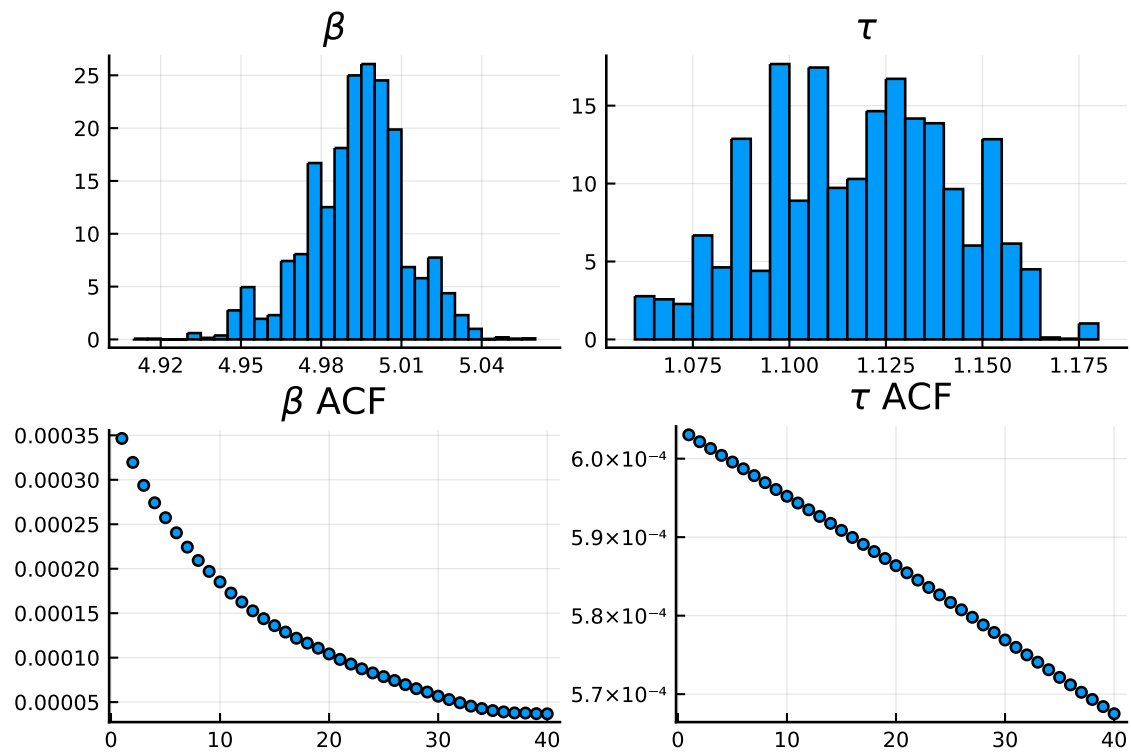


Fig. 2: Metropolis-Hastings Linear Link

on some form of Gaussian quadrature or importance sampling, there could be a much better mixture distribution for the taus.

For testing the Logistic link function, I sampled the distribution of the Y by using $Y_{i,t} \sim \text{bernoulli}(g(X_{i,t}\beta + \alpha_i))$ where g is the inverse of the logit function. This is the standard link function for logistic regression. Conditional on β, α the likelihood function is then given by:

$$L = \prod_{i=1}^I \prod_{t=1}^T y_{i,t} \log(g(X_{i,t}\beta + \alpha_i)) + (1 - y_{i,t}) \log(1 - g(X_{i,t}\beta + \alpha_i))$$

Using Algorithm A, I find a much better mixing probability than the under the linear model, which is likely because I do not simulate the model under any noise, and only round up or down based on the inverse-logit transformation. I find an acceptance probability of 41%. The distribution for τ is quite different, with much of its mass to the left of the mode rather than a symmetric distribution that was expected.

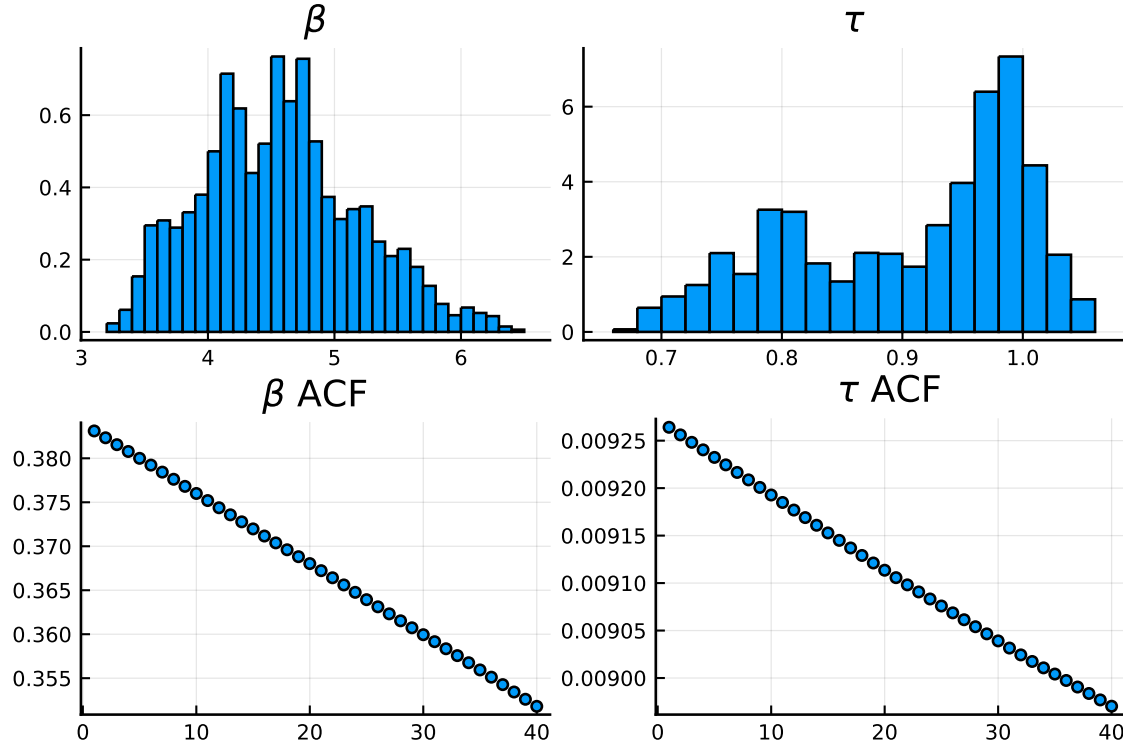


Fig. 3: Logistic Algorithm A

For the pseudo-marginal MCMC sampler of the distribution in the logistic, the distribution is far less ideal. The estimate for β has most of its mass away from the true value that the simulated data is based around. While the β estimate is underestimated, the τ is overestimated to attempt to compensate. I find that there is an extremely high acceptance rate of 95%, which is up quite a lot from the linear case, though the correlation is much higher as well. This indicates that there is not as much exploration of the distribution, and that it is not mixing as well as it should.

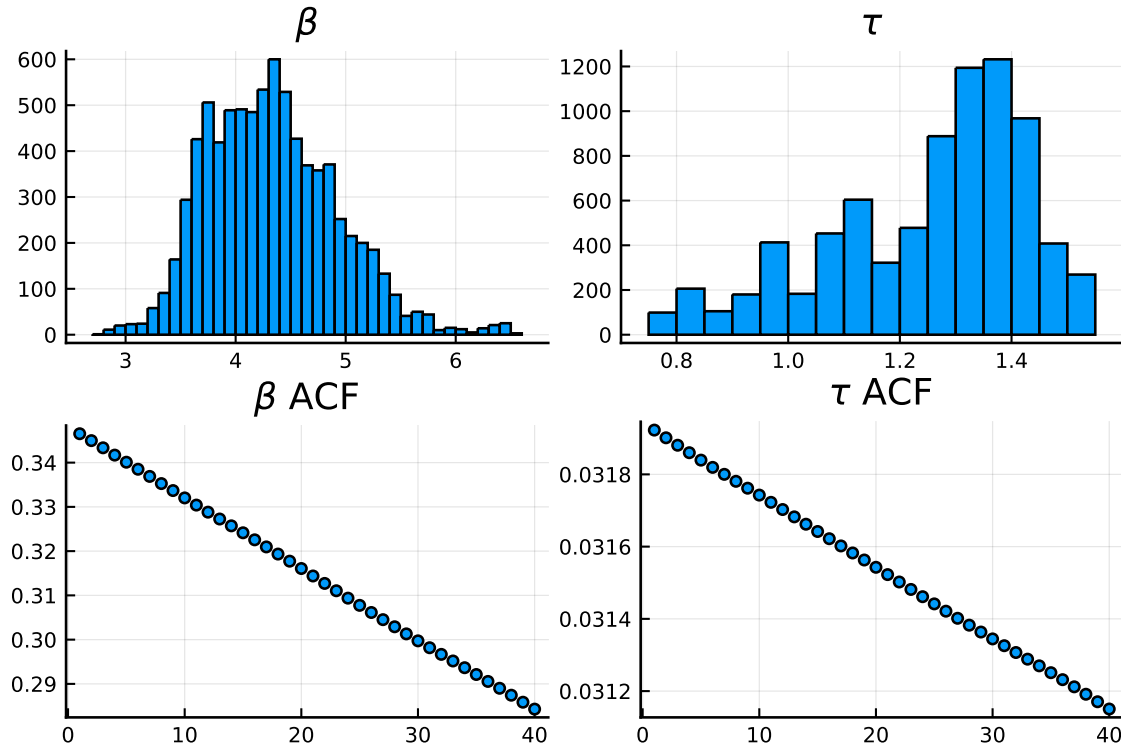


Fig. 4: Logistic Pseudo-Marginal MCMC

In this case, we see that the MCMC actually performs better than the pseudo-marginal MCMC sampler for this distribution. This is for a relatively large sample, and a large size for the MCMC as well as a large number of samples for α .

If we consider the linear case, and wish to figure out the distribution of the likelihood function, we may calculate the likelihood values for the sampled values of β, τ .